

1 (a)

```
1. function log_prior(zs)
2.     return factorized_gaussian_log_density(0,0,zs)
3. end
```

1 (b)

```
1. function logp_a_beats_b(za,zb)
2.     return -log1pexp.(-(za-zb))
3. end
```

1 (c)


```
1. function all_games_log_likelihood(zs,games)
2.     zs_a = zs[games[:,1],:]
3.     zs_b = zs[games[:,2],:]
4.     likelihoods = logp_a_beats_b.(zs_a,zs_b)
5.     return sum(likelihoods, dims=1)
6. end
```

1 (d)


```
1. function joint_log_density(zs,games)
2.     return sum(all_games_log_likelihood(zs,games).+log_prior(zs),dims=1)
3. end
4. @testset "Test shapes of batches for likelihoods" begin
5.     B = 15 # number of elements in batch
6.     N = 4 # Total Number of Players
7.     test_zs = randn(4,15)
8.     test_games = [1 2; 3 1; 4 2] # 1 beat 2, 3 beat 1, 4 beat 2
9.     @test size(test_zs) == (N,B)
10.    #batch of priors
11.    @test size(log_prior(test_zs)) == (1,B)
12.    # loglikelihood of p1 beat p2 for first sample in batch
13.    @test size(logp_a_beats_b(test_zs[1,1],test_zs[2,1])) == ()
14.    # loglikelihood of p1 beat p2 broadcasted over whole batch
15.    @test size(logp_a_beats_b.(test_zs[1,:],test_zs[2,:])) == (B,)
16.    # batch loglikelihood for evidence
17.    @test size(all_games_log_likelihood(test_zs,test_games)) == (1,B)
18.    # batch loglikelihood under joint of evidence and prior
19.    @test size(joint_log_density(test_zs,test_games)) == (1,B)
20. end
```

Test Summary:	Pass	Total
Test shapes of batches for likelihoods	6	6


2 (a)



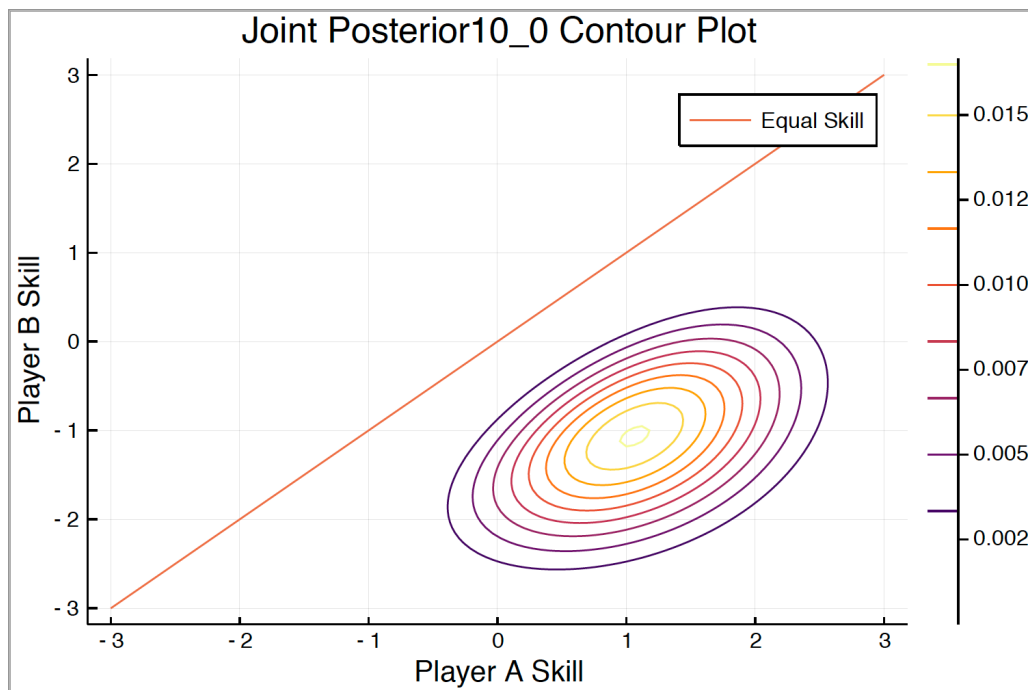
2 (b)



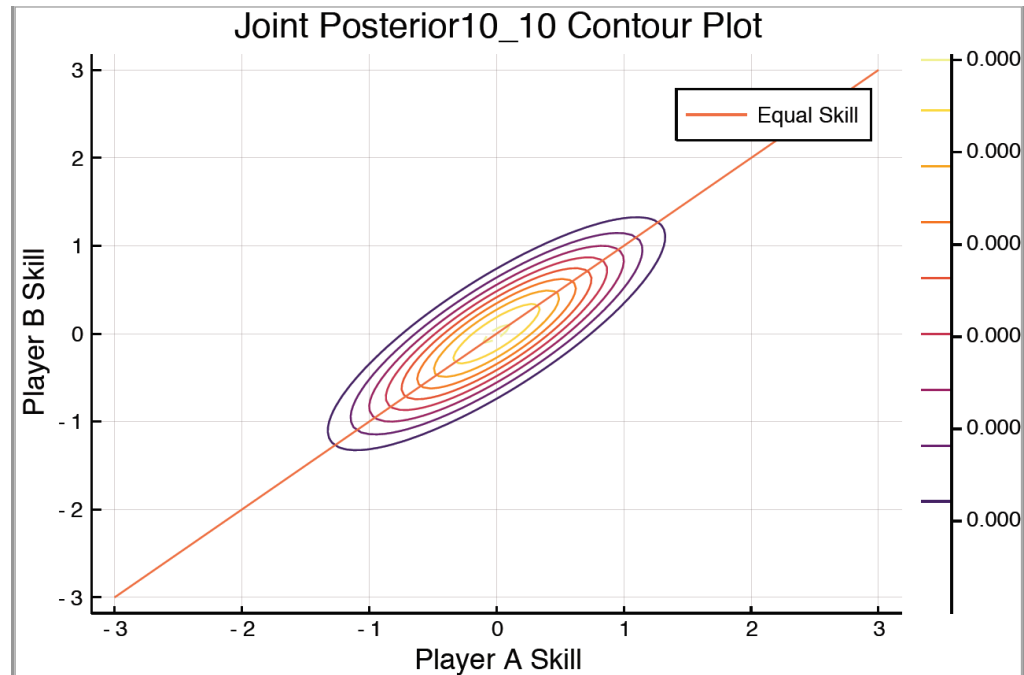
2 (c)



2 (d)



2 (e)



3 (a)

```
1. function elbo(params,logp,num_samples)
2.   samples = exp.(params[2]) .*randn(length(params[1]),num_samples) .+params[1]
3.   logp_estimate = logp(samples)
4.   logq_estimate = factorized_gaussian_log_density(params[1],params[2],samples)
5.   return sum(logp_estimate-logq_estimate)/num_samples
6. end
```

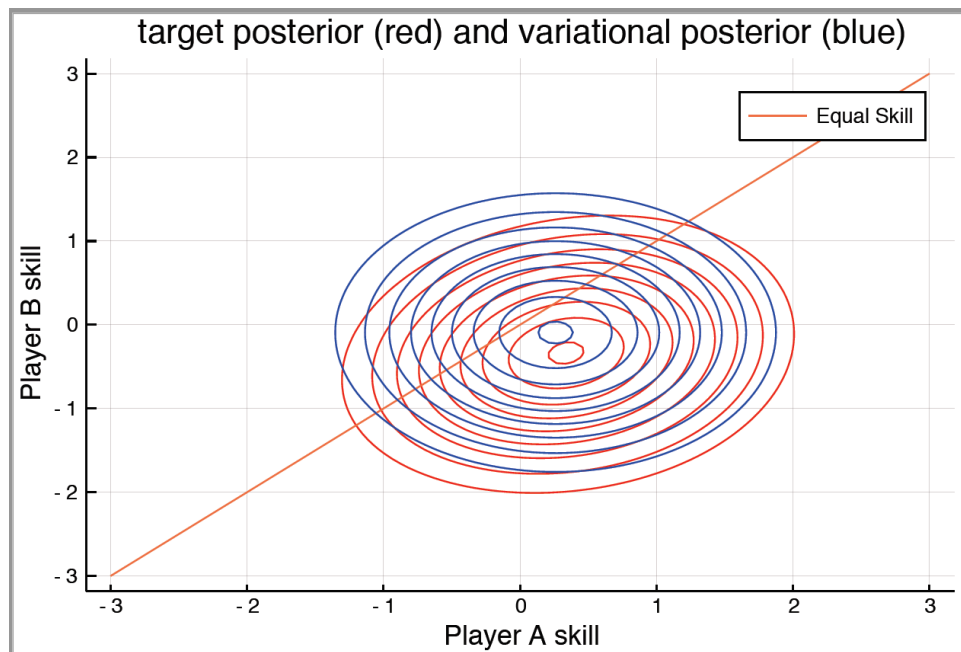
3 (b)

```
1. function neg_toy_elbo(params; games = two_player_toy_games(1,0), num_samples = 10
   0)
2.   logp(zs) = joint_log_density(zs,games)
3.   return -elbo(params,logp, num_samples)
4. end
```


3 (c)

```
1. function fit_toy_variational_dist(init_params, toy_evidence; num_itr=200, lr= 1e-
2, num_q_samples = 10)
2.   params_cur = init_params
3.   for i in 1:num_itr
4.     grad_params = gradient(params -> neg_toy_elbo(params; games = toy_evidence, num_
samples = num_q_samples), params_cur)[1] #gradients of variational objective with respec
t to parameters
5.     params_cur = params_cur .- lr.* grad_params #update paramters with lr-
sized step in descending gradient
6.     @info "nelbo: $(neg_toy_elbo(params_cur; games=toy_evidence, num_samples=num
_q_samples))" #report the current elbbo during training
7.     # plot true posterior in red and variational in blue
8.     plot();
9.     #plot likelihood contours for target posterior
10.    display(skillcontour!(zs->exp(joint_log_density(zs, toy_evidence)),colour=:red))
11.    plot_line_equal_skill!()
12.    #plot likelihood contours for variational posterior
13.    display(skillcontour!(zs -> exp(factorized_gaussian_log_density(params_cur[1], param
s_cur[2],zs)),colour=:blue))
14.  end
15.  return params_cur
16. end
```


3 (d)



3 (e)



3 (f)



4 (a)

Question: Are the isocontours of  $P(z_i, z_j | \text{all games})$  the same as those of  $P(z_i, z_j | \text{games between } i \text{ and } j)$ ?

Answer: Yes

Question: do the games between other players besides  $i$  and  $j$  provide information about the skill of players  $i$  and  $j$ ?


Answer: No

4(b)

```
1. function fit_variational_dist(init_params, tennis_games;
2. num_itr=200, lr= 1e-2, num_q_samples = 10)
3.     params_cur = init_params
4.     for i in 1:num_itr
5.         grad_params = gradient(params -> neg_toy_elbo(params; games = tennis_games, num_
        samples = num_q_samples), params_cur)[1]
6.         params_cur = params_cur .- lr.*grad_params
7.         @info "nelbo: $(neg_toy_elbo(params_cur; games=tennis_games, num_samples=num
        _q_samples))" #report objective value with current parameters
8.     end
9.     return params_cur
10. end
11. init_mu = randn(107) # random initialization
12. init_log_sigma = randn(107) # random initialization
13. init_params = (init_mu, init_log_sigma)
14. trained_params = fit_variational_dist(init_params, tennis_games)
```

The final negative ELBO estimate after optimization is 1144.024976806034.

4 (c)



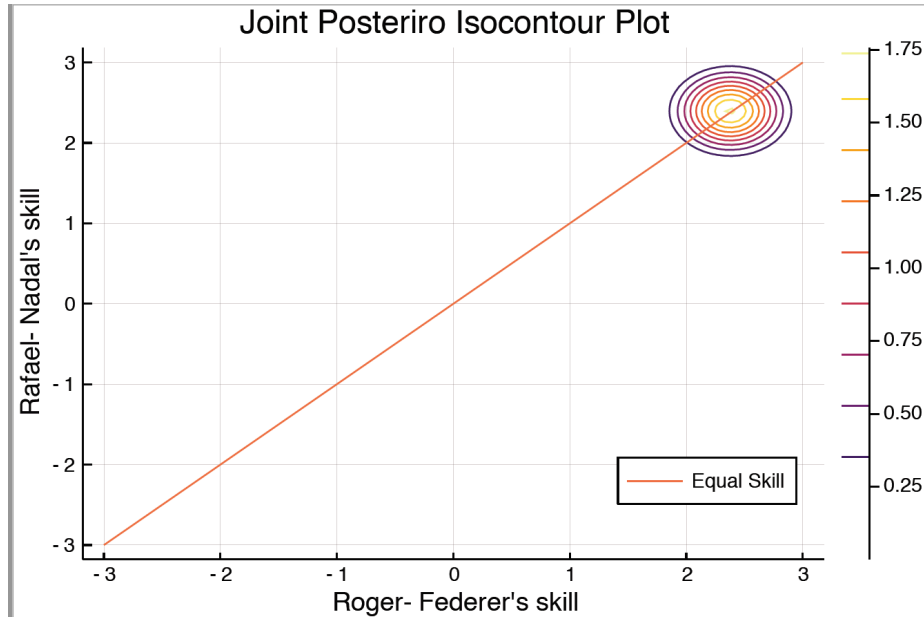
4 (d)

The top 10 players ranked by the highest mean skills are:

Novak-Djokovic  
Roger-Federer  
Rafael-Nadal  
Andy-Murray  
Robin-Soderling  
David-Ferrer  
Jo-Wilfried-Tsonga  
Tomas-Berdych  
Juan-Martin-Del-Potro  
Richard-Gasquet



4 (e)



4 (f)

Since  $y_a = z_a - z_b$  and  $y_b = z_b$

We can write  $Y = AZ$  where  $Y = \begin{pmatrix} y_a \\ y_b \end{pmatrix}$ ,  $A = \begin{pmatrix} 1 & -1 \\ 0 & 1 \end{pmatrix}$  and  $Z = \begin{pmatrix} z_a \\ z_b \end{pmatrix}$

By hint 2, since  $Z \sim N(\mu_z, \Sigma_z)$  where  $\mu_z = \begin{pmatrix} \mu_a \\ \mu_b \end{pmatrix}$  and  $\Sigma_z = \begin{pmatrix} \sigma_{aa} & 0 \\ 0 & \sigma_{bb} \end{pmatrix}$ , we can get  $Y \sim N(A\mu_z, A\Sigma_z A^T)$ .

$$A\mu_z = \begin{pmatrix} \mu_a - \mu_b \\ \mu_b \end{pmatrix} \text{ and } A\Sigma_z A^T = \begin{pmatrix} 1 & -1 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} \sigma_{aa} & 0 \\ 0 & \sigma_{bb} \end{pmatrix} \begin{pmatrix} 1 & 0 \\ -1 & 1 \end{pmatrix} = \begin{pmatrix} \sigma_{aa} + \sigma_{bb} & -\sigma_{bb} \\ -\sigma_{bb} & \sigma_{bb} \end{pmatrix}$$

By hint 3, we can get  $y_a \sim N(\mu_a - \mu_b, \sigma_{aa} + \sigma_{bb})$

We want to know  $P(z_a > z_b) \equiv P(z_a - z_b > 0) \equiv P(y_a > 0)$

We need to use cdf function to get the probability.

1-cdf(Normal(mean, sigma),0)

4 (g)

The exact probability under the approximate posterior that Roger Federer has higher skill than Rafael Nadal is 0.5421.

Calculated by the following code:

```
1. # P(zRF>zRN)
2. print(1-cdf(Normal(means[5]-means[1], sqrt((exp(logsd[1]))^2+(exp(logsd[5]))^2)),0))
```

The estimated probability by simple Monte Carlo with 10000 examples is 0.5448.

Calculated by the following code:

```
1. # P(zRF>zRN)
2. zs_RF = rand(Normal(means[5], exp(logsd[5])),10000)
3. zs_RN = rand(Normal(means[1], exp(logsd[1])),10000)
4. print(sum(zs_RF .> zs_RN,dims=1) / 10000)
```

4 (h)

The exact probability under the approximate posterior that Roger Federer is better than the player with the lowest mean skill is  $0.999 \approx 1$ .

Calculated by the following code:

```
1. # P(zRF>z_b)
2. print(1-cdf(Normal(means[5]-means[75], sqrt((exp(logsd[75]))^2+(exp(logsd[5]))^2)),0))
```

The estimated probability by simple Monte Carlo with 10000 examples is 1.

Calculated by the following code:

```
1. # P(zRF>z_b)
2. lowestindex = perm[107]
3. zs_RF = rand(Normal(means[5], exp(logsd[5])),10000)
4. zs_b = rand(Normal(means[75], exp(logsd[75])),10000)
5. print(sum(zs_RF .> zs_b,dims=1) / 10000)
```

4 (i)

If we change the prior to  $\text{Normal}(10,1)$ , then the posterior will be changed based on the prior. Since “elbo” function is designed based on the prior then “neg\_toy\_elbo” function and “fit\_variational\_dist” function will change as well. So all questions related with above functions will change.

Therefore, Q4 b, c, e, g will change.

For Q4 b, the info neg elbo changed if we change the prior distribution.

For Q4 c, the mean of all players will increase by approximate 10.

For Q4 e, the plot will shift to upper right corner.

For Q4 g, the probability increases from 0.54 to 0.59.

For Q4 h, the probability is still 1 since the difference between the lowest player and RF is too large. Hard to see the change.

## Appendix

The following code is for plotting the above plots.

```
1. # Convenience function for producing toy games between two players.
2. two_player_toy_games(p1_wins, p2_wins) = vcat([repeat([1,2]',p1_wins), repeat([2,1]',p2
   _wins)]...)
3.
4. # Example for how to use contour plotting code
5. #plot(title="Example Gaussian Contour Plot",
6. #      xlabel = "Player 1 Skill",
7. #      ylabel = "Player 2 Skill"
8. #      )
9. #example_gaussian(zs) = exp(factorized_gaussian_log_density([-1.,2.],[0.,0.5],zs))
10. #skillcontour!(example_gaussian)
11. #plot_line_equal_skill!()
12. #savefig(joinpath("plots","example_gaussian.pdf"))
13.
14.
15. zs = randn(2,15)
16. # plot prior contours
17. plot(title="Prior Contour Plot",
18.       xlabel = "Player A Skill",
19.       ylabel = "Player B Skill"
20.       )
21. prior(zs) = exp(log_prior(zs))
22. skillcontour!(zs -> prior(zs))
23. plot_line_equal_skill!()
24. savefig(joinpath("plots","joint_prior.pdf"))
25.
26. # plot likelihood contours
```

```

27. game12 = two_player_toy_games(1,2)
28. plot(title="Likelihood Contour Plot",
29.       xlabel = "Player A Skill",
30.       ylabel = "Player B Skill"
31.       )
32. likelihood(zs) = exp(all_games_log_likelihood(zs,game12))
33. skillcontour!(zs -> likelihood(zs))
34. plot_line_equal_skill!()
35. savefig(joinpath("plots","likelihood.pdf"))
36.
37. # plot joint contours with player A winning 1 game
38. games = two_player_toy_games(1,0)
39. plot(title="Joint Posterior1_0 Contour Plot",
40.       xlabel = "Player A Skill",
41.       ylabel = "Player B Skill"
42.       )
43. joint1(zs) = exp(joint_log_density(zs,games))
44. skillcontour!(zs -> joint1(zs))
45. plot_line_equal_skill!()
46. savefig(joinpath("plots","joint1.pdf"))
47. # plot joint contours with player A winning 10 games
48. games = two_player_toy_games(10,0)
49. plot(title="Joint Posterior10_0 Contour Plot",
50.       xlabel = "Player A Skill",
51.       ylabel = "Player B Skill"
52.       )
53. joint10(zs) = exp(joint_log_density(zs,games))
54. skillcontour!(zs -> joint10(zs))
55. plot_line_equal_skill!()
56. savefig(joinpath("plots","joint10.pdf"))
57.
58. # plot joint contours with player A winning 10 games and player B winning 10 games
59. games = two_player_toy_games(10,10)
60. plot(title="Joint Posterior10_10 Contour Plot",
61.       xlabel = "Player A Skill",
62.       ylabel = "Player B Skill"
63.       )
64. joint20(zs) = exp(joint_log_density(zs,games))
65. skillcontour!(zs -> joint20(zs))
66. plot_line_equal_skill!()
67. savefig(joinpath("plots","joint20.pdf"))
68.
69. #fit q with SVI observing player A winning 1 game
70. game10 = two_player_toy_games(1,0)
71. fit_toy_variational_dist(toy_params_init,game10)
72. xlabel!("Player A skill")
73. ylabel!("Player B skill")
74. title!("target posterior (red) and variational posterior (blue)")
75. savefig(joinpath("plots","posterior 1.pdf"))
76.
77. #fit q with SVI observing player A winning 10 games
78. game100 = two_player_toy_games(10,0)
79. fit_toy_variational_dist(toy_params_init,game100)
80. xlabel!("Player A skill")
81. ylabel!("Player B skill")
82. title!("target posterior (red) and variational posterior (blue)")
83. savefig(joinpath("plots","posterior 10.pdf"))
84.
85. #fit q with SVI observing player A winning 10 games and player B winning 10 games
86. game1010 = two_player_toy_games(10,10)
87. fit_toy_variational_dist(toy_params_init,game1010)

```

```

88. xlabel!("Player A skill")
89. ylabel!("Player B skill")
90. title!("target posterior (red) and variational posterior (blue)")
91. savefig(joinpath("plots", "posterior 20.pdf"))
92.
93. #10 players with highest mean skill under variational model
94. #hint: use sortperm
95. means = trained_params[1]
96. logsd = trained_params[2]
97. perm = sortperm(means, rev=true)
98. plot(means[perm], yerror = exp.(logsd[perm]))
99. xlabel!("Sorted Player by skills")
100. ylabel!("players' skills")
101. title!("Approx. mean (line) and variance (vertical range) sorted by skills")
102. savefig(joinpath("plots", "mean_variance under variational model.pdf"))
103.
104. #joint posterior over "Roger-Federer" and "Rafael-Nadal"
105. #hint: findall function to find the index of these players in player_names
106. indexRF = findall(name-> name == "Roger-Federer", player_names)
107. indexRN = findall(name-> name == "Rafael-Nadal", player_names)
108. plot(legend=:bottomright)
109. skillcontour!(zs -> exp(factorized_gaussian_log_density([means[1], means[5]], [logsd[1], logsd[5]], zs)))
110. plot_line_equal_skill!()
111. xlabel!("Roger-Federer skill")
112. ylabel!("Rafael-Nadal skill")
113. title!("Joint Posterior Isocontour Plot")
114. savefig(joinpath("plots", "Joint Posterior Isocontour Plot.pdf"))

```