

1. (a)

The code to compute the log of the prior is

```
log_prior(z) = factorized_gaussian_log_density(0,0,z)
```

1. (b)

The code to produce a 784-dimensional mean vector of a product of Bernoulli distributions is

```
Dz, Dh = 2, 500  
Ddata = 28^2  
decoder = Chain(Dense(Dz,Dh,tanh), Dense(Dh, Ddata))
```

1. (c)

The code to compute the log-likelihood  $\log p(x|z)$  is

```
function log_likelihood(x,z)  
  logitmean = decoder(z)  
  return sum(bernoulli_log_density(logitmean,x) , dims = 1)  
end
```

1. (d)

The code to compute the joint log density  $\log p(z, x)$  is

```
joint_log_density(x,z) = log_prior(z) .+ log_likelihood(x,z)
```

2. (a)

The code to output the mean and log-standard deviation of a factorized Gaussian is

```
encoder = Chain(Dense(Ddata,Dh,tanh),Dense(Dh,2*Dz),unpack_gaussian_params)
```

2. (b)

The code to evaluate the likelihood of  $z$  under the variational distribution is

```
log_q(q_μ, q_logσ, z) = factorized_gaussian_log_density(q_μ, q_logσ,z)
```

2. (c)

The code to compute an unbiased estimate of the mean variational evidence lower bound on a batch of images is

```

function elbo(x)
    q_μ = encoder(x)[1]
    q_logσ = encoder(x)[2]
    z = exp.(q_logσ) .* randn(size(q_μ)) .+ q_μ
    joint_ll = joint_log_density(x,z)
    log_q_z = log_q (q_μ, q_logσ, z)
    elbo_estimate = sum(joint_ll.-log_q_z) / size(q_μ)[2]
    return elbo_estimate
end

```

2. (d)

The code to compute the negative elbo estimate over a batch of data is

```

function loss(x)
    return -elbo(x)
end

```

2. (e)

The code to initialize and optimize the encoder and decoder parameters jointly on the training set is

```

function train_model_params!(loss, encoder, decoder, train_x, test_x; nepochs=10)
    ps = Flux.params(encoder,decoder)
    opt = ADAM()
    for i in 1:nepochs
        for d in batch_x(train_x)
            gs = Flux.gradient(ps) do
                batch_loss = loss(d)
                return batch_loss
            end
            Flux.Optimise.update!(opt,ps,gs)
        end
        if i%1 == 0 # change 1 to higher number to compute and print less frequently
            @info "Test loss at epoch $i: $(loss(batch_x(test_x)[1]))"
        end
    end
    @info "Parameters of encoder and decoder trained!"
end

```

## Train the model

```
train_model_params!(loss,encoder,decoder,train_x,test_x, nepochs=100)
```

We train the data for 100 epochs and find the final ELBO on the test set is 150.64.

3. (a)

The following code generates the 2x10 plots . The first row plots are the Bernoulli means of  $p(x|z)$  and the second row plots are the binary images, sampled from the distribution above it.


```
z = randn(2,10)
logitp = decoder(z)
p = (1 .+ exp.(-logitp)) .^ (-1)
x = sample_bernoulli(p)
plot_list_mean = []
for i in 1:10
    push!(plot_list_mean, plot(mnist_img(θ[:,i])))
end
plot_list_images = []
for j in 1:10
    push!(plot_list_images, plot(mnist_img(x[:,j])))
end
plot_list = [plot_list_mean; plot_list_images]
display(plot(plot_list..., layout = grid(2,10), size = (10000,6000)))
```



Generative Sample Plots (1<sup>st</sup> Row: mean, 2<sup>nd</sup> Row: binarized)

3. (b)

```
q_μ, q_logσ = encoder(train_x)[1], encoder(train_x)[2]  
scatter(q_μ[1:], q_μ[2:], group = train_label)
```



3. (c)

In the figure below, each row represents the plots of the corresponding pairs.


```
function interpolation(za,zb, $\alpha$ )
    return  $\alpha$  * za .+ (1- $\alpha$ ) * zb
end
# Manually pick the data with different classes and check by train_label
train_label[21:30] # 4, 0, 9, 1, 1, 2, 4, 3
x1 = train_x[:,21] # Number 4
x2 = train_x[:,22] # Number 0
x3 = train_x[:,23] # Number 9
x4 = train_x[:,24] # Number 1
x5 = train_x[:,26] # Number 2
x6 = train_x[:,30] # Number 3

logitmean1 = encoder(x1)[1]
logitmean2 = encoder(x2)[1]
logitmean3 = encoder(x3)[1]
logitmean4 = encoder(x4)[1]
logitmean5 = encoder(x5)[1]
logitmean6 = encoder(x6)[1]

bernmean1 = []
for i in 0:10
    logitmeans = decoder(interpolation(logitmean1, logitmean2, i/10))
    means = exp.(logitmeans) ./ (1 .+ exp.(logitmeans))
    push!(bernmean1, plot(mnist_img(means[:,1])))
end

bernmean2 = []
for i in 0:10
    logitmeans = decoder(interpolation(logitmean3, logitmean4, i/10))
    means = exp.(logitmeans) ./ (1 .+ exp.(logitmeans))
    push!(bernmean2, plot(mnist_img(means[:,1])))
end

bernmean3 = []
for i in 0:10
    logitmeans = decoder(interpolation(logitmean5, logitmean6, i / 10))
    means = exp.(logitmeans) ./ (1 .+ exp.(logitmeans))
    push!(bernmean3, plot(mnist_img(means[:,1])))
end
bernmean = []
bernmean = [bernmean1; bernmean2; bernmean3]
display(plot(bernmean..., layout=grid(3,11), size=(10000, 2000), axis=nothing))
```



Bernoulli Means from the Latent Space Interpolation

4. (a) (a)

```
function top_image(x)
    " return only the top half of a 28*28 array"
end
```

4. (a) (b)

```
function log_likelihood_top(x,z)
    logitmean = decoder(z)
    logitmean_top # gives the logit mean of the top half images
    sum(bernoulli_log_density(logitmean_top, x), dims=1)
end
```

4. (a) (c)

```
log_joint_top(x,z) = log_prior(z) .+ log_likelihood_top(x,z)
```

4. (b)

```
encoder_2 = Chain(Dense(392,Dh, tanh), Dense(Dh,4), unpack_gaussian_params())
```

```
function elbo_2(x)
    q_μ = encoder_2(x)[1]
    q_logσ = encoder_2(x)[2]
    z = exp.(q_logσ) .* randn(size(q_μ)) .+ q_μ
    joint_ll = log_joint_top(x,z)
    log_q_z = factorized_gaussian_log_density(q_μ, q_logσ, z)
    elbo_estimate_2 = sum(joint_ll.-log_q_z) / size(q_μ)[2]
    return elbo_estimate_2
end
```

```

function loss_2(x)
    return -elbo_2(x)
end

```

```

function train_model_params_2!(loss, encoder, decoder, train_x, test_x; nepochs=10)
    ps = Flux.params(encoder,decoder)
    opt = ADAM()
    for i in 1:nepochs
        for d in batch_x(train_x)
            gs = Flux.gradient(ps) do
                batch_loss = loss(d)
                return batch_loss
            end
            Flux.Optimise.update!(opt,ps,gs)
        end
        if i%1 == 0 # change 1 to higher number to compute and print less frequently
            @info "Test loss at epoch $i: $(loss(batch_x(test_x)[1]))"
        end
    end
    @info "Parameters of encoder and decoder trained!"
end end
@info "Parameters of encoder and decoder trained!"
end

train_model_params_2!(loss_2,encoder_2,decoder,train_x, test_x, nepochs
=100)

```

4. (c)
- (a) True
  - (b) False
  - (c) False
  - (d) False
  - (e) True